

Autoencoder anomaly detection

October 2, 2023

```
[37]: import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, \
    precision_score

RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]
```

```
[38]: dataset = pd.read_csv("F:\College Data\Deep L\PV 4\DL\creditcard.csv")
```

```
[39]: #check for any null values
print("Any nulls in the dataset",dataset.isnull().values.any())
print('-----')
print("No. of unique labels",len(dataset['Class'].unique()))
print("Label values",dataset.Class.unique())


#0 is for normal credit card transcation
#1 is for fraudulent credit card transcation
print('-----')
print("Break down of Normal and Fraud Transcations")
print(pd.value_counts(dataset['Class'],sort=True))
```

```
Any nulls in the dataset False
-----
No. of unique labels 2
Label values [0 1]
-----
Break down of Normal and Fraud Transcations
Class
0      284315
1         492
```

Name: count, dtype: int64

```
[40]: #visualizing the imbalanced dataset
count_classes = pd.value_counts(dataset['Class'],sort=True)
count_classes.plot(kind='bar',rot=0)
plt.xticks(range(len(dataset['Class'].unique())),dataset.Class.unique())
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
```


[40]: Text(0, 0.5, 'Number of Observations')



```
[41]: #Save the normal and fraudulent transactions in separate dataframe
normal_dataset = dataset[dataset.Class == 0]
fraud_dataset = dataset[dataset.Class == 1]

#Visualize transaction amounts for normal and fraudulent transactions
bins = np.linspace(200,2500,100)
plt.hist(normal_dataset.Amount,bins=bins,alpha=1,density=True,label='Normal')
plt.hist(fraud_dataset.Amount,bins=bins,alpha=0.5,density=True,label='Fraud')
```

```
plt.legend(loc='upper right')
plt.title("Transcation Amount vs Percentage of Transcations")
plt.xlabel("Transcation Amount (USD)")
plt.ylabel("Percentage of Transcations")
plt.show()
```



```
[42]: dataset
```

```
[42]:
```

	Time	V1	V2	V3	V4	V5	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	
...	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	

284806 172792.0 -0.533413 -0.189733 0.703337 -0.506271 -0.012546

	V6	V7	V8	V9	...	V21	V22	\
0	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	
1	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	
2	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	
3	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	
4	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	
...	
284802	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	
284803	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	
284804	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	
284805	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	
284806	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	

	V23	V24	V25	V26	V27	V28	Amount	\
0	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	
1	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	
2	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	
3	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	
4	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	
...	
284802	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	
284803	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	
284804	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	
284805	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	
284806	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	

	Class
0	0
1	0
2	0
3	0
4	0
...	...
284802	0
284803	0
284804	0
284805	0
284806	0

[284807 rows x 31 columns]

```
[43]: sc = StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))
```

```
[44]: raw_data = dataset.values
      #The last element contains if the transaction is normal which is represented by 0
      #and if fraud then 1
      labels = raw_data[:, -1]

      #The other data points are the electrocardiogram data
      data = raw_data[:, 0:-1]

      train_data, test_data, train_labels, test_labels = train_test_split(
          data, labels, test_size = 0.2, random_state = 2021)
```

```
[45]: min_val = tf.reduce_min(train_data)
      max_val = tf.reduce_max(train_data)

      train_data = (train_data - min_val) / (max_val - min_val)
      test_data = (test_data - min_val) / (max_val - min_val)

      train_data = tf.cast(train_data, tf.float32)
      test_data = tf.cast(test_data, tf.float32)
```

```
[46]: train_labels = train_labels.astype(bool)
      test_labels = test_labels.astype(bool)

      #Creating normal and fraud datasets
      normal_train_data = train_data[~train_labels]
      normal_test_data = test_data[~test_labels]

      fraud_train_data = train_data[train_labels]
      fraud_test_data = test_data[test_labels]
      print("No. of records in Fraud Train Data=", len(fraud_train_data))
      print("No. of records in Normal Train Data=", len(normal_train_data))
      print("No. of records in Fraud Test Data=", len(fraud_test_data))
      print("No. of records in Normal Test Data=", len(normal_test_data))
```

```
No. of records in Fraud Train Data= 389
No. of records in Normal Train Data= 227456
No. of records in Fraud Test Data= 103
No. of records in Normal Test Data= 56859
```

```
[47]: nb_epoch = 20
      batch_size = 64
      input_dim = normal_train_data.shape[1]
      #num of columns, 30
      encoding_dim = 14
      hidden_dim1 = int(encoding_dim / 2)
      hidden_dim2 = 4
      learning_rate = 1e-7
```

```
[48]: #input layer
input_layer = tf.keras.layers.Input(shape=(input_dim,))

#Encoder
encoder = tf.keras.layers.
    ↳Dense(encoding_dim,activation="tanh",activity_regularizer = tf.keras.
    ↳regularizers.l2(learning_rate))(input_layer)
encoder = tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim2,activation=tf.nn.
    ↳leaky_relu)(encoder)

#Decoder
decoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
decoder = tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim,activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim,activation='tanh')(decoder)

#Autoencoder
autoencoder = tf.keras.Model(inputs = input_layer,outputs = decoder)
autoencoder.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 30)]	0
dense_12 (Dense)	(None, 14)	434
dropout_4 (Dropout)	(None, 14)	0
dense_13 (Dense)	(None, 7)	105
dense_14 (Dense)	(None, 4)	32
dense_15 (Dense)	(None, 7)	35
dropout_5 (Dropout)	(None, 7)	0
dense_16 (Dense)	(None, 14)	112
dense_17 (Dense)	(None, 30)	450

```
=====  
Total params: 1168 (4.56 KB)  
Trainable params: 1168 (4.56 KB)  
Non-trainable params: 0 (0.00 Byte)
```

```

[49]: cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.
      ↪h5",mode='min',monitor='val_loss',verbose=2,save_best_only=True)
      #Define our early stopping
      early_stop = tf.keras.callbacks.EarlyStopping(
          monitor='val_loss',
          min_delta=0.0001,
          patience=10,
          verbose=11,
          mode='min',
          restore_best_weights=True
      )

[50]: autoencoder.compile(metrics=['accuracy'],loss=_,
      ↪'mean_squared_error',optimizer='adam')

[51]: history = autoencoder.fit(normal_train_data,normal_train_data,epochs = nb_epoch,
      batch_size = batch_size,shuffle = True,
      validation_data = (test_data,test_data),
      verbose=1,
      callbacks = [cp,early_stop]).history

```

Epoch 1/20

3545/3554 [=====>.] - ETA: 0s - loss: 0.0035 - accuracy: 0.0305

Epoch 1: val_loss improved from inf to 0.00003, saving model to autoencoder_fraud.h5

C:\Users\Acer\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\engine\training.py:3000: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
 saving_api.save_model(

3554/3554 [=====] - 14s 3ms/step - loss: 0.0035 - accuracy: 0.0304 - val_loss: 2.5545e-05 - val_accuracy: 0.0024

Epoch 2/20

3541/3554 [=====>.] - ETA: 0s - loss: 1.9480e-05 - accuracy: 0.0564

Epoch 2: val_loss improved from 0.00003 to 0.00002, saving model to autoencoder_fraud.h5

3554/3554 [=====] - 11s 3ms/step - loss: 1.9475e-05 - accuracy: 0.0563 - val_loss: 2.3230e-05 - val_accuracy: 0.0024

Epoch 3/20

3548/3554 [=====>.] - ETA: 0s - loss: 1.9425e-05 - accuracy: 0.0657

Epoch 3: val_loss improved from 0.00002 to 0.00002, saving model to

```

autoencoder_fraud.h5
3554/3554 [=====] - 12s 3ms/step - loss: 1.9421e-05 -
accuracy: 0.0658 - val_loss: 2.2306e-05 - val_accuracy: 0.0010
Epoch 4/20
3547/3554 [=====>.] - ETA: 0s - loss: 1.9538e-05 -
accuracy: 0.0587
Epoch 4: val_loss improved from 0.00002 to 0.00002, saving model to
autoencoder_fraud.h5
3554/3554 [=====] - 12s 3ms/step - loss: 1.9538e-05 -
accuracy: 0.0586 - val_loss: 2.0339e-05 - val_accuracy: 0.0024
Epoch 5/20
3553/3554 [=====>.] - ETA: 0s - loss: 1.9482e-05 -
accuracy: 0.0615
Epoch 5: val_loss did not improve from 0.00002
3554/3554 [=====] - 11s 3ms/step - loss: 1.9483e-05 -
accuracy: 0.0615 - val_loss: 2.0421e-05 - val_accuracy: 0.0661
Epoch 6/20
3552/3554 [=====>.] - ETA: 0s - loss: 1.9032e-05 -
accuracy: 0.0696
Epoch 6: val_loss did not improve from 0.00002
3554/3554 [=====] - 11s 3ms/step - loss: 1.9030e-05 -
accuracy: 0.0695 - val_loss: 2.1059e-05 - val_accuracy: 0.0514
Epoch 7/20
3551/3554 [=====>.] - ETA: 0s - loss: 1.8035e-05 -
accuracy: 0.1084
Epoch 7: val_loss did not improve from 0.00002
3554/3554 [=====] - 11s 3ms/step - loss: 1.8035e-05 -
accuracy: 0.1084 - val_loss: 2.3646e-05 - val_accuracy: 0.0257
Epoch 8/20
3538/3554 [=====>.] - ETA: 0s - loss: 1.7710e-05 -
accuracy: 0.1286
Epoch 8: val_loss did not improve from 0.00002
3554/3554 [=====] - 12s 3ms/step - loss: 1.7702e-05 -
accuracy: 0.1286 - val_loss: 2.5736e-05 - val_accuracy: 0.0257
Epoch 9/20
3541/3554 [=====>.] - ETA: 0s - loss: 1.7467e-05 -
accuracy: 0.1511
Epoch 9: val_loss did not improve from 0.00002
3554/3554 [=====] - 12s 3ms/step - loss: 1.7465e-05 -
accuracy: 0.1511 - val_loss: 2.5645e-05 - val_accuracy: 0.0270
Epoch 10/20
3552/3554 [=====>.] - ETA: 0s - loss: 1.7224e-05 -
accuracy: 0.1891
Epoch 10: val_loss did not improve from 0.00002
3554/3554 [=====] - 12s 3ms/step - loss: 1.7226e-05 -
accuracy: 0.1891 - val_loss: 2.7170e-05 - val_accuracy: 0.0284
Epoch 11/20
3548/3554 [=====>.] - ETA: 0s - loss: 1.6969e-05 -

```




```
accuracy: 0.2355
Epoch 11: val_loss did not improve from 0.00002
Restoring model weights from the end of the best epoch: 1.
3554/3554 [=====] - 12s 3ms/step - loss: 1.6970e-05 -
accuracy: 0.2356 - val_loss: 2.4134e-05 - val_accuracy: 0.0262
Epoch 11: early stopping
```

```
[52]: plt.plot(history['loss'],linewidth = 2,label = 'Train')
plt.plot(history['val_loss'],linewidth = 2,label = 'Test')
plt.legend(loc='upper right')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')

#plt.ylim(ymin=0.70,ymax=1)

plt.show()
```




```
[53]: test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2),axis = 1)
```

```
error_df = pd.DataFrame({'Reconstruction_error':mse,
                        'True_class':test_labels})
```

1781/1781 [=====] - 3s 2ms/step

```
[54]: threshold_fixed = 50
groups = error_df.groupby('True_class')
fig,ax = plt.subplots()


for name,group in groups:
    ax.plot(group.index,group.Reconstruction_error,marker='o',ms = 3.
    ↪5,linestyle='',
            label = "Fraud" if name==1 else "Normal")
ax.hlines(threshold_fixed,ax.get_xlim()[0],ax.
    ↪get_xlim()[1],colors="r",zorder=100,label="Threshold")
ax.legend()
plt.title("Reconstructions error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show()
```



```
[55]: threshold_fixed = 52
pred_y = [1 if e > threshold_fixed else 0
          for e in
            error_df.Reconstruction_error.values]
error_df['pred'] = pred_y
conf_matrix = confusion_matrix(error_df.True_class,pred_y)

plt.figure(figsize = (4,4))
sns.heatmap(conf_matrix,xticklabels = LABELS,yticklabels = LABELS,annot = 
    ↪True,fmt="d")
plt.title("Confusion matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()

#Print Accuracy,Precision and Recall
print("Accuracy :",accuracy_score(error_df['True_class'],error_df['pred']))
print("Recall :",recall_score(error_df['True_class'],error_df['pred']))
print("Precision :",precision_score(error_df['True_class'],error_df['pred']))
```



Accuracy : 0.9981917769741231
 Recall : 0.0
 Precision : 0.0

```
C:\Users\Acer\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 due to no predicted samples. Use  
`zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))
```

```
[ ]:
```

```
[ ]:
```